

# **Data Definition Language (DDL), Views and Indexes**

**Instructor: Shel Finkelstein**

*Reference:*

*A First Course in Database Systems,  
3<sup>rd</sup> edition, Chapter 2.3 and 8.1-8.4*

# Important Notices

- Lab3 assignment is due on **Sunday, May 19**, by 11:59pm.
  - 3 weeks to do Lab3 because of Midterm.
  - Lab2 was/will be discussed at Lab Sections.
  - Your solution should be submitted via Canvas as a zip file.
    - Canvas is used for both Lab submission and grading.
    - Late Lab Assignments will not be accepted.
    - Be sure that you post the correct file!
  - Load file for Lab3 has been posted to Piazza.
    - You must use load file to do Lab3.
    - Helps with testing, but we won't post query solutions.
- The **third** Gradiance Assignment, "CMPS 182 Spring 2019 #3", was/will be posted on Monday, April 29, and is due **Monday, May 6**, by 11:59pm.

# Important Notices

- **Reminder:** Midterm is on **Wednesday, May 8.**
  - **Includes material up to and including previous Lecture**, but not this Lecture.
  - No make-ups, no early Midterms, no late Midterms ... and no devices.
  - You may bring a **single two-sided 8.5" x 11" sheet of paper** with as much info written (or printed) on it as you can fit and read unassisted.
    - No sharing of these sheets will be permitted.
  - “Practice Midterm” from Spring 2017 has been posted on Piazza under Resources → Exams.
    - Solution will be posted there next week ... but take it yourself first.
    - Some questions are on topics we haven't covered yet.
  - Hope that all requests for DRC accommodation have been submitted.
  - Piazza announcement will describe required seating pattern for Midterm.
- See [Small Group Tutoring website](#) for LSS Tutoring with [Chandler Hawkins](#).

# Important Notices **Midterm**

- **Reminder:** Midterm is on **Monday, Feb 11.**
  - **Includes material up to and including Lecture 6**, but not this Lecture, Lecture 7.
  - No make-ups, no early Midterms, no late Midterms ... and no devices.
  - You may bring a **single two-sided 8.5" x 11" sheet of paper** with as much info written (or printed) on it as you can fit and read unassisted.
    - No sharing of these sheets will be permitted.
  - “Practice Midterm” from Fall 2018 has been posted on Piazza.
    - Solution will be posted on Friday, February 8 ... but take it yourself first.
  - Hope that requests for DRC accommodation have all been submitted.
    - *You should have already received notification of DRC room if you submitted your DRC request!*
  - Piazza announcement describes required seating pattern for Midterm.
    - You'll be assigned alternating seats with one space between them as you enter the room, and then gaps will be filled as assigned by the TAs.

# SQL Language

- Data Manipulation Language (DML)
  - Access and modify data
  - SELECT, INSERT, DELETE, UPDATE
- Data Definition Language (DDL)
  - Modify structure of data
  - CREATE, DROP, ALTER
- Data Control Language (DCL)
  - Control access to the data (security)
  - GRANT, REVOKE
- Databases also have Utilities, such as Backup/Restore
  - Syntax not specified in the SQL standard

# CREATE TABLE

```
CREATE TABLE MovieStar (  
    name            CHAR(30) ,  
    address         VARCHAR(255) DEFAULT 'Hollywood',  
    gender          CHAR(1),  
    birthdate       DATE NOT NULL DEFAULT '2001-12-30',  
    PRIMARY KEY (name)  
);
```

- PRIMARY KEY
- DEFAULT
- NOT NULL

# Reminder: Some Facts About Nulls

- Almost all comparisons with NULL will evaluate to unknown. If Salary is NULL, then the following will be UNKNOWN:
  - Salary = 10
  - Salary <> 10
  - 90 > Salary OR 90 <= Salary
  - Salary = NULL
  - Salary <> NULL
- Use of IS NULL and IS NOT NULL
  - Salary IS NULL will be true if Salary is NULL, false otherwise
  - Salary IS NOT NULL will be true if Salary isn't NULL, false otherwise
- ORDER BY works with attributes that can have NULL values
  - NULL will probably be smallest or largest value
  - Not specified by SQL standard, so it depends on the implementation
- GROUP BY also works with attributes that can have NULL values

# DROP TABLE

- Dropping a table:

```
DROP TABLE MovieStar;
```

- Don't assume that rolling back transaction will bring back the table!
  - Interaction of DDL and transactions may depend on implementation.



# ALTER TABLE

- Adding a column to a table:
  - ALTER TABLE MovieStar ADD phone CHAR(16) DEFAULT 'unlisted';
- Dropping a column from a table:
  - ALTER TABLE MovieStar DROP birthdate;
  - In some systems:  
ALTER TABLE MovieStar DROP **COLUMN** birthdate;
  - In some SQL systems, dropping a column isn't allowed.
- Changing the type of a column:
  - Some implementations let you change type of column in limited ways.

# What Can You CREATE/DROP in SQL DDL?

- TABLE
- **VIEW**
- **INDEX**
- ASSERTION
- TRIGGER
- SCHEMA
- PROCEDURE/FUNCTION/TYPE
  - SQL2003 standard, but there are significant variations in implementations in different systems
- ...

# VIEWS: Motivation for Views

- Views help with **logical data independence**, allowing you to retrieve data if it matches the description in the view.

```
CREATE VIEW < view-name> AS < view-definition> ;
```

```
CREATE VIEW ParamountMovies AS  
    SELECT title, year  
    FROM Movies  
    WHERE studioName = 'Paramount ' ;
```

- You may now ask queries on ParamountMovies **as if it were a table**:  
 SELECT title FROM ParamountMovies WHERE year=1976;  
 – Composition in SQL is powerful: Tables, Queries, Views

# Some Advantages of Views

- **Short-hand/encapsulation:** You can treat a view as a table (for queries), which allows you to define a short-hand for a concept that might involved a complicated query.
- **Re-use:** You can re-use a view as often as you like; other people who can access the view may also re-use it.
- **Authorization:** People may be granted access to a view, even though they don't have access to the underlying tables.
  - They may not even know that the view isn't a table.
  - Even if they know that it's a view, they may not know which tables underlie it.
- **Logical data independence:** Even if the tables underlying a view change, the view may still be used in queries (or applications), without re-writing the queries (or applications).

# More Views

Movies ( title , year , length , genre , studioName , producerC# )

MovieExec ( name , address , cert# , netWorth )

```
CREATE VIEW MovieProd AS
```

```
    SELECT m.title, e.name, m.genre
```

```
    FROM Movies m, MovieExec e
```

```
    WHERE m.producerC# = e.cert# ;
```

```
SELECT DISTINCT genre
```

```
FROM MovieProd
```

```
WHERE name = 'George Lucas';
```

# Renaming Attributes in CREATE VIEW

Movies ( title , year , length , genre , studioName , producerC# )

MovieExec ( name , address , cert# , netWorth )

```
CREATE VIEW MovieProd(movie_title, prod_name, movie_genre) AS
```

```
    SELECT m.title, e.name, m.genre
```

```
    FROM Movies m , MovieExec e
```

```
    WHERE m.producerC# = e.cert# ;
```

```
SELECT DISTINCT movie_genre
```

```
FROM MovieProd
```

```
WHERE prod_name = 'George Lucas';
```

# What is a View?

- A view can include any SQL SELECT statement
  - Including UNION, Aggregates, GROUP BY, HAVING, ORDER BY, etc.
- A view is not stored as a table
  - The tables underlying the view are stored in the database, but only the description of the view is in the database
  - ... although some systems support MATERIALIZED VIEWS
- But a view can be used in many (not all) of the same ways as tables
  - Views can be queried
  - Views can be defined on views, as well as on tables!

# Queries on Views and Tables

```
CREATE VIEW ParamountMovies AS
  SELECT title , year
  FROM Movies
  WHERE studioName = 'Paramount ' ;
```

```
SELECT DISTINCT s.starName
FROM ParamountMovies p , StarsIn s
WHERE p.title = s.movieTitle AND p.year = s.movieYear ;
```

```
CREATE VIEW ParamountStars AS
  SELECT DISTINCT starName
  FROM ParamountMovies , StarsIn
  WHERE title = movieTitle AND year = movieYear ;
```



# DROP VIEW

```
CREATE VIEW ParamountMovies AS  
    SELECT title , year  
    FROM Movies  
    WHERE studioName = 'Paramount ' ;
```

```
DROP View ParamountMovies;
```

- What happens if you execute the following after dropping that view?
  - SELECT \* FROM ParamountMovies;
  - SELECT \* FROM Movies;

# View Updates

- Some modification operations on views work, but others do not, generally failing either because:
  - Constraint on underlying table would be violated, or
  - The effects of the View modification is not well-defined on the underlying tables.
- This is a complex topic, which we'll only discuss briefly.
  - See Textbook Section 8.2 for more info.

# View Update Problems

Movies(title, year, length, genre, studioName , producerC#)

```
CREATE VIEW ParamountMovies AS
```

```
  SELECT title , year
```

```
  FROM Movies
```

```
  WHERE studioName = 'Paramount ' ;
```

```
INSERT INTO ParamountMovies VALUES ('StarTrek', 1979);
```

The INSERT will fail if the other columns of Movies (besides title and year) don't have defaults, and also don't allow NULL values.

# **View Update Problems (continued)**

Ambiguous View Update example with  
Employees and Departments

<< We'll draw this on the board >>

# INDEXES: Motivation for Indexes

- Searching an entire table may take a long time:

```
SELECT *
```

```
FROM Movies
```

```
WHERE studioName = 'Disney' AND year = 1990;
```

If there were 100 Million movies, searching them might take a while. An index (e.g., a B-Tree) would allow faster access to matching movies.

If a table is updated, all Indexes on that table are immediately automatically updated within the same transaction.

- Which indexes do you need to change on INSERT and DELETE?
- What about UPDATE?

# CREATE INDEX

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney' AND year = 1990;
```

How much would each of these indexes help?

```
CREATE INDEX YearIndex ON Movies(year);
```

```
CREATE INDEX StudioIndex ON Movies(studioName);
```

```
CREATE INDEX YSIndex ON Movies(year,studioName);
```

```
CREATE INDEX SYIndex ON Movies(studioName,year);
```

How much would each of the indexes help if the WHERE clause was just  
**year = 1990?**

# Indexes and Ordering

```
SELECT *  
FROM Movies  
WHERE studioName = 'Disney' AND year < 1990;
```

How much would each of these indexes help?

```
CREATE INDEX YearIndex ON Movies(year);
```

```
CREATE INDEX StudioIndex ON Movies(studioName);
```

```
CREATE INDEX YSIndex ON Movies(year,studioName);
```

```
CREATE INDEX SYIndex ON Movies(studioName,year);
```

How much would each of the indexes help if the WHERE clause was just **year < 1990**?

# Indexes and Physical Independence

- SQL statements can be executed regardless of which indexes (if any) exist in the database.
  - Applications don't have to be modified when indexes are created or dropped!
- What gets impacted when indexes are created or dropped?
  - Performance of SQL statements
  - Some may run faster, some may run slower



# Disadvantages of Indexes?

- Why not put indexes on every attributes, or even on every combination of attributes that you might query on?
  - Huge number of indexes
  - Space for indexes
  - Cache impact of searching indexes
  - Update time for indexes when table is modified

# Index Design

- Most Database Administrators (DBAs) pick a set of indexes that work well on expected workload, and there are tools that help pick good indexes
  - But workloads change, so choice of indexes may need to change
    - `DROP INDEX YearIndex;`
- Keys are indexed (automatically in many database systems) to:
  - Help maintain uniqueness (primary key, unique)
  - Check Foreign Key references to Primary Keys (Referential Integrity)

# Index Utilization

- SQL statements don't specify use of indexes, so they don't have to be modified when you change what's indexed!
  - Database Optimizer tries to figure out "the best"/"a good" way to execute each SQL query.
  - All the tuples in a Relation can be scanned directly, without using indexes, so indexes aren't necessary ... except for performance.
  - Some systems have ways that you can tell the Optimizer what to do. This has advantages and disadvantages. (What are they?)
- Many SQL systems (including PostgreSQL) have an EXPLAIN PLAN statement, so that you can see what plan the optimizer chooses for a SQL statement.